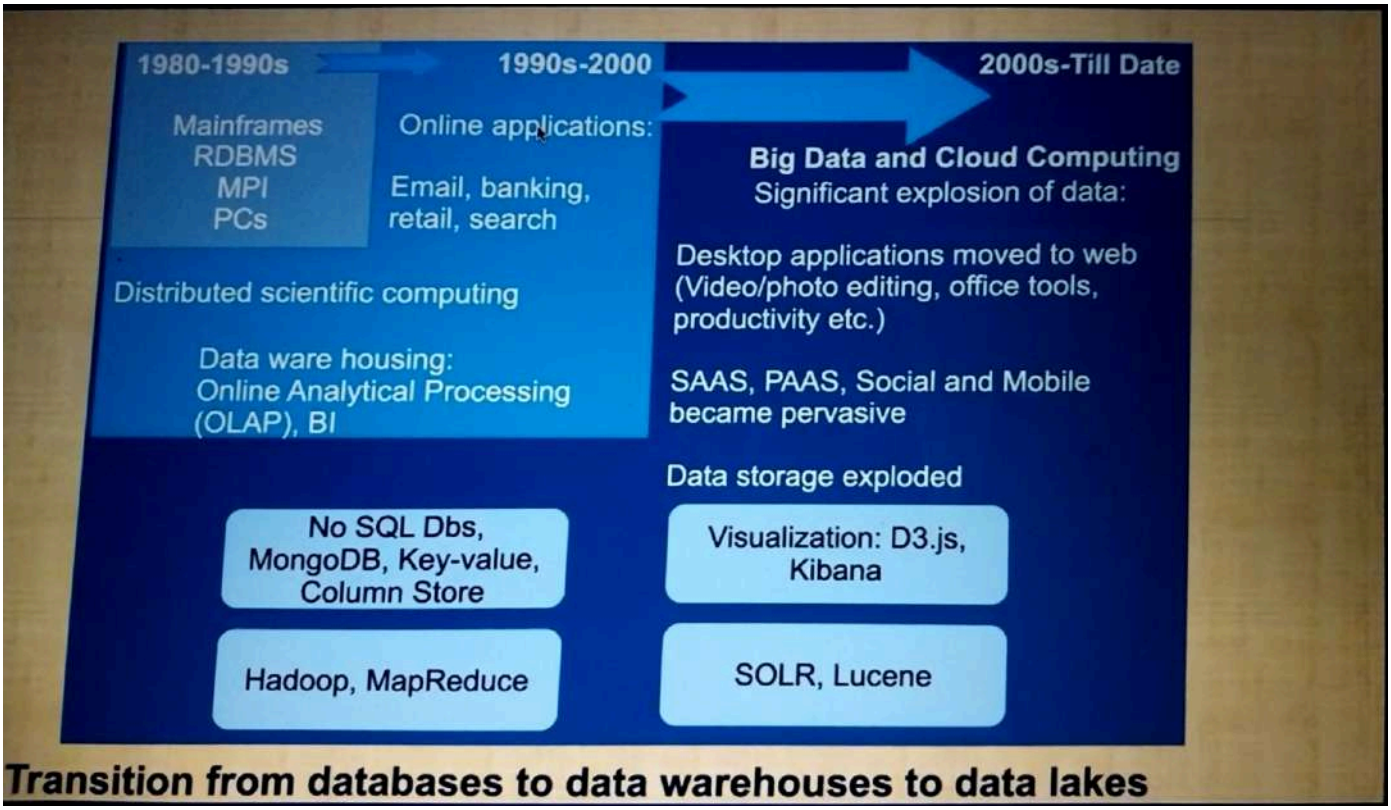

Transition from databases to data warehouses to data lakes

* Curriculum Discussion :- * Big Data - Hadoop & spark
* Cloud - Azure & GCP
- Devop Tools - Git, Github & Docker
- No.SQL - mongoDB. Tool

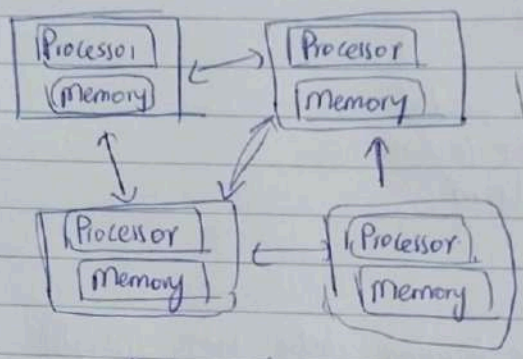

Distributed Computing

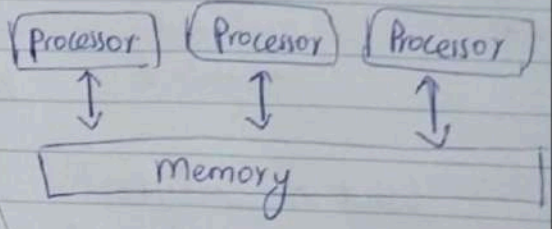

Parallel - Computing

Cluster Computing

- A cluster is a type of parallel or distributed processing sy which consist of a collection of interconnected stand-alone computers cooperatively working together as a single

Big data : It ~~too~~ Data that is Too large & too complex for conventional data tools to capture, store & analyze

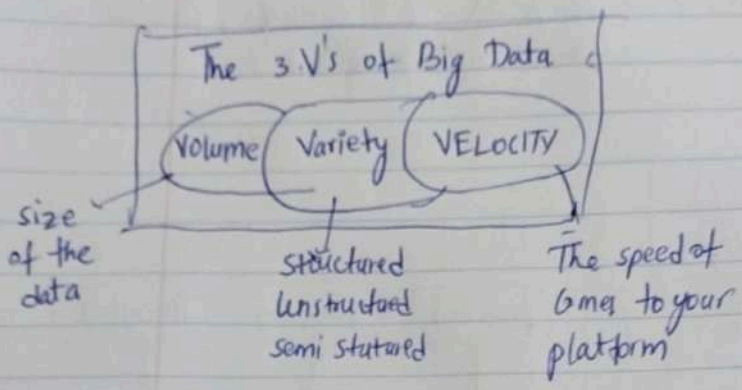

What is analytics?

The scientific process of transforming data into insight for making better decisions, offering new opportunities for a competative advantage.

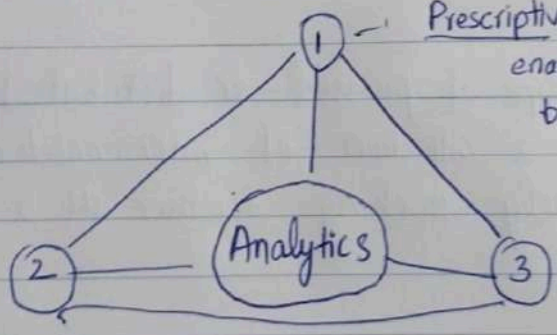

Prescriptive Analytics
enabling smart decisions based on data.
What should we do?

Predictive analytics
Predicting the future based on historical patterns
What could happen?

Descriptive analytics
Mining data to provide business insights
What has happened?

Why do airlane prices change every hour?

Ans - Prescriptive analytics
(advice on possible outcomes).

How do grocery cashiers know to hand you coupons you might actually use?

Ans Predictive Analytics
(understanding the future).

How does Netflix frequently recommend just the right movie?

Descriptive Analytics
(insight into the past).

Problems with Traditional approach

1) Storing huge and exponentially growing datasets.
2). Processing data having complex structure (structured, un-structured, -semi-structured)
3) Bringing huge amount of data to computation unit becoma a bottlened

ieo1548083867.mp4

Rack awareness

Replica storage is a tradeoff between **reliability** and read/write **bandwidth**.

To increase reliability, we need to store block replicas on different racks and Datanodes to increase fault tolerance. While the write bandwidth is lowest when replicas are stored on the same node. Therefore, Hadoop has a default strategy to deal with this **conundrum**, also known as the **Rack Awareness algorithm**.

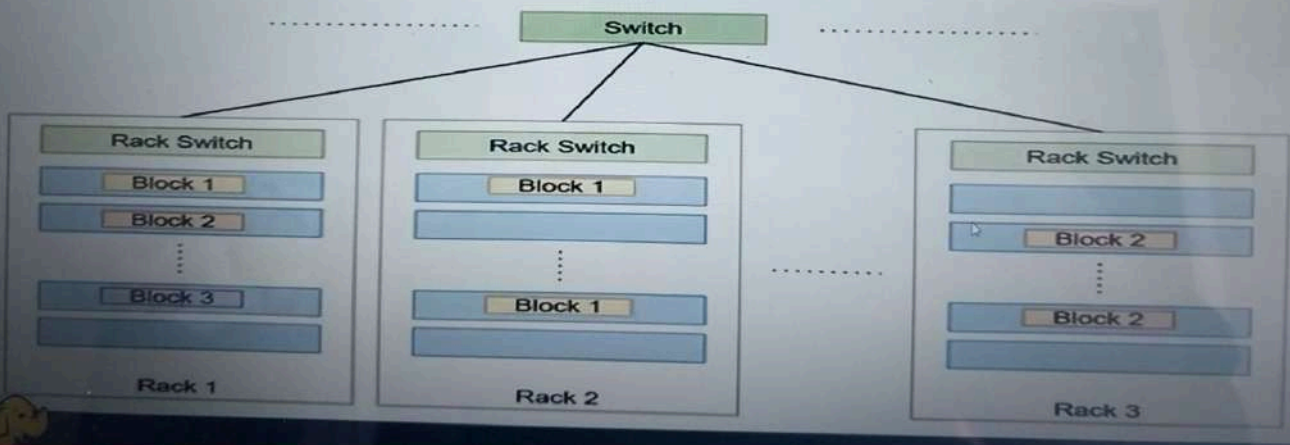

1:05:33 / 1:33:38

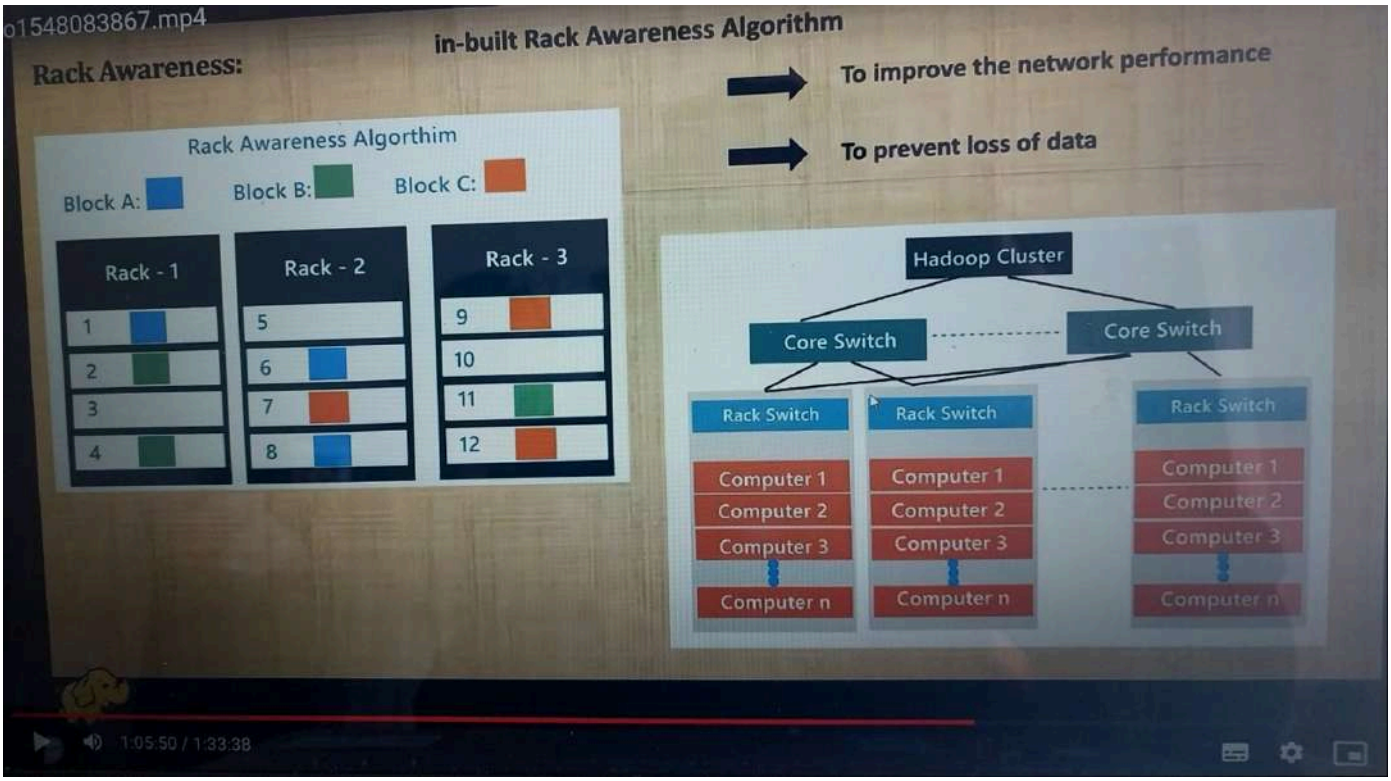
01548083867.mp4
Rack Awareness:
in-built Rack Awareness Algorithm
To improve the network performance
To prevent loss of data
Rack Awareness Algorthim
Block A:
Block B:
Block C:
Rack - 1
Rack - 2
Rack - 3
1
2
3
4
5
6
7
8
9
10
11
12
Hadoop Cluster
Core Switch
Core Switch
Rack Switch
Rack Switch
Rack Switch
Computer 1
Computer 2
Computer 3
Computer n
Computer 1
Computer 2
Computer 3
Computer n
Computer 1
Computer 2
Computer 3
Computer n
1:05:50 / 1:33:38

(The data storing copies) → by default it has 3 copies

2). Replication → TA -1 Hadoop → 3 copies

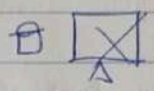

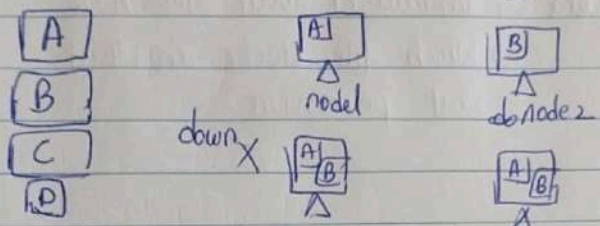

→ How replication works)

Rack awarness Algorithm → Replication → Data Failure, Network Performance

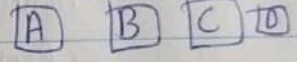

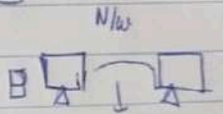

↳ The performance is high ✓ but it we lost Rack1 we can lose total data so it will store atmost 2 Rack

Rack Awarnell algo → atmost -2 Racks.

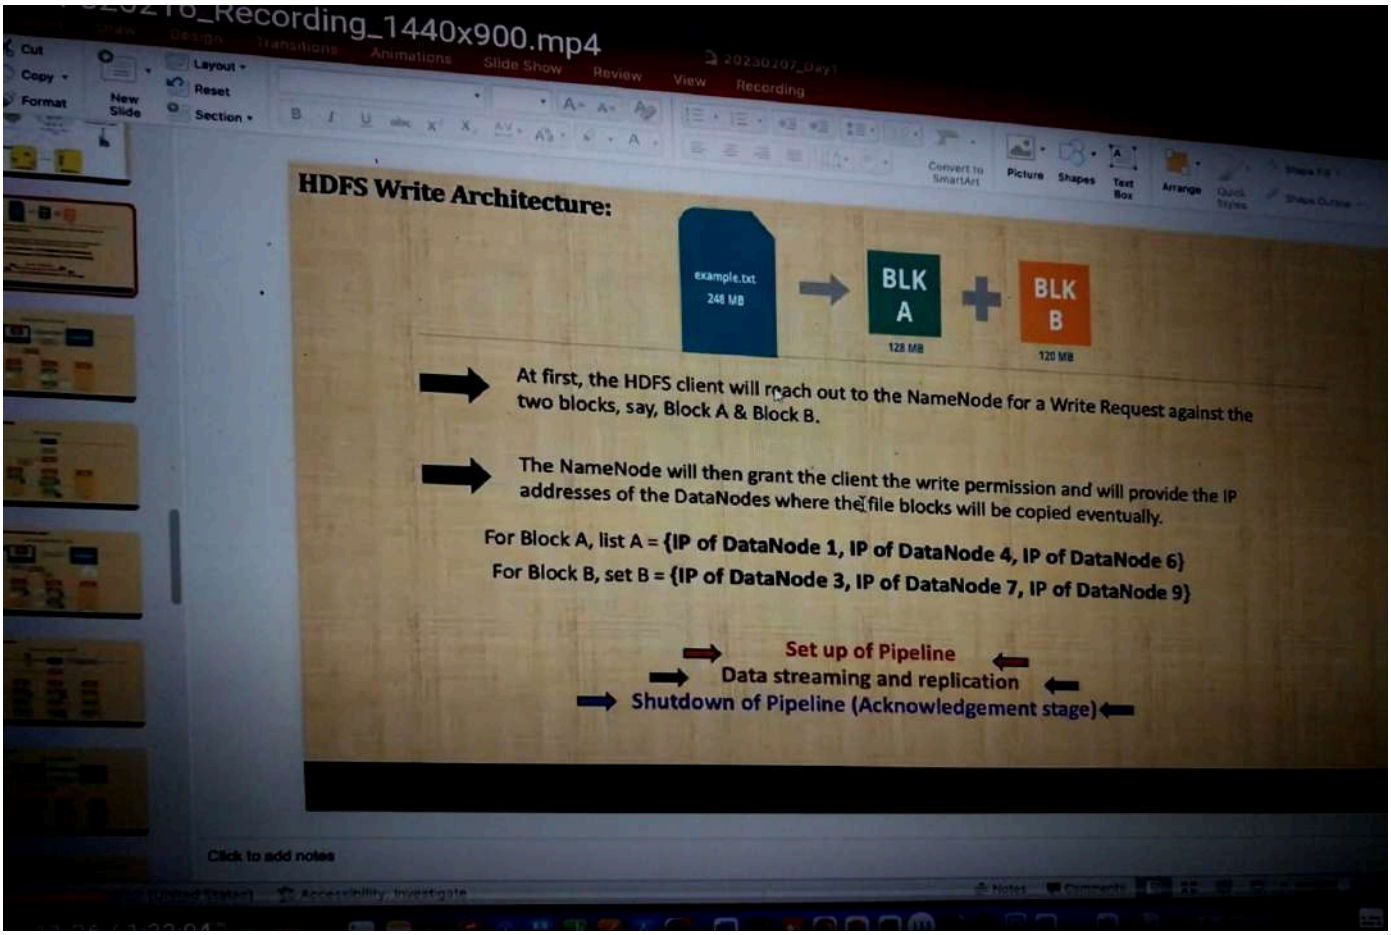
_Recording_1440x900.mp4
HDFS Write Architecture:
example.txt
248 MB
BLK A
128 MB
BLK B
120 MB
At first, the HDFS client will reach out to the NameNode for a Write Request against the two blocks, say, Block A & Block B.
The NameNode will then grant the client the write permission and will provide the IP addresses of the DataNodes where the file blocks will be copied eventually.
For Block A, list A = {IP of DataNode 1, IP of DataNode 4, IP of DataNode 6}
For Block B, set B = {IP of DataNode 3, IP of DataNode 7, IP of DataNode 9}
Set up of Pipeline
Data streaming and replication
Shutdown of Pipeline (Acknowledgement stage)
Click to add notes

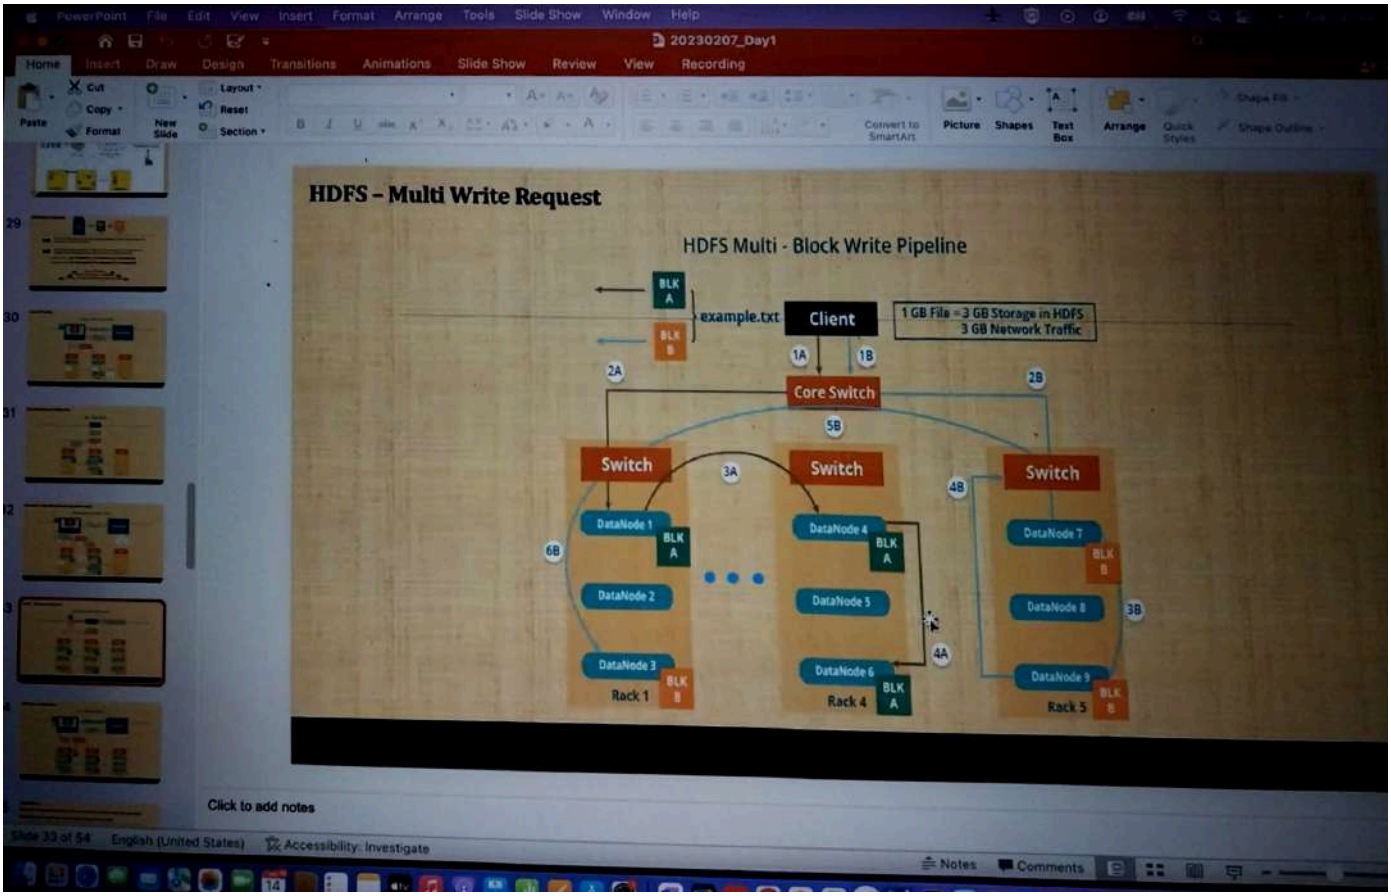
HDFS – Multi Write Request
HDFS Multi - Block Write Pipeline
example.txt
BLK A
BLK B
Client
1 GB File = 3 GB Storage in HDFS
3 GB Network Traffic
1A
1B
2A
2B
Core Switch
5B
Switch
3A
Switch
4B
Switch
6B
DataNode 1
BLK A
DataNode 2
DataNode 3
BLK B
Rack 1
DataNode 4
BLK A
DataNode 5
DataNode 6
BLK A
Rack 4
DataNode 7
BLK B
DataNode 8
3B
4A
DataNode 9
BLK B
Rack 5
Click to add notes

HDFS Read Architecture:

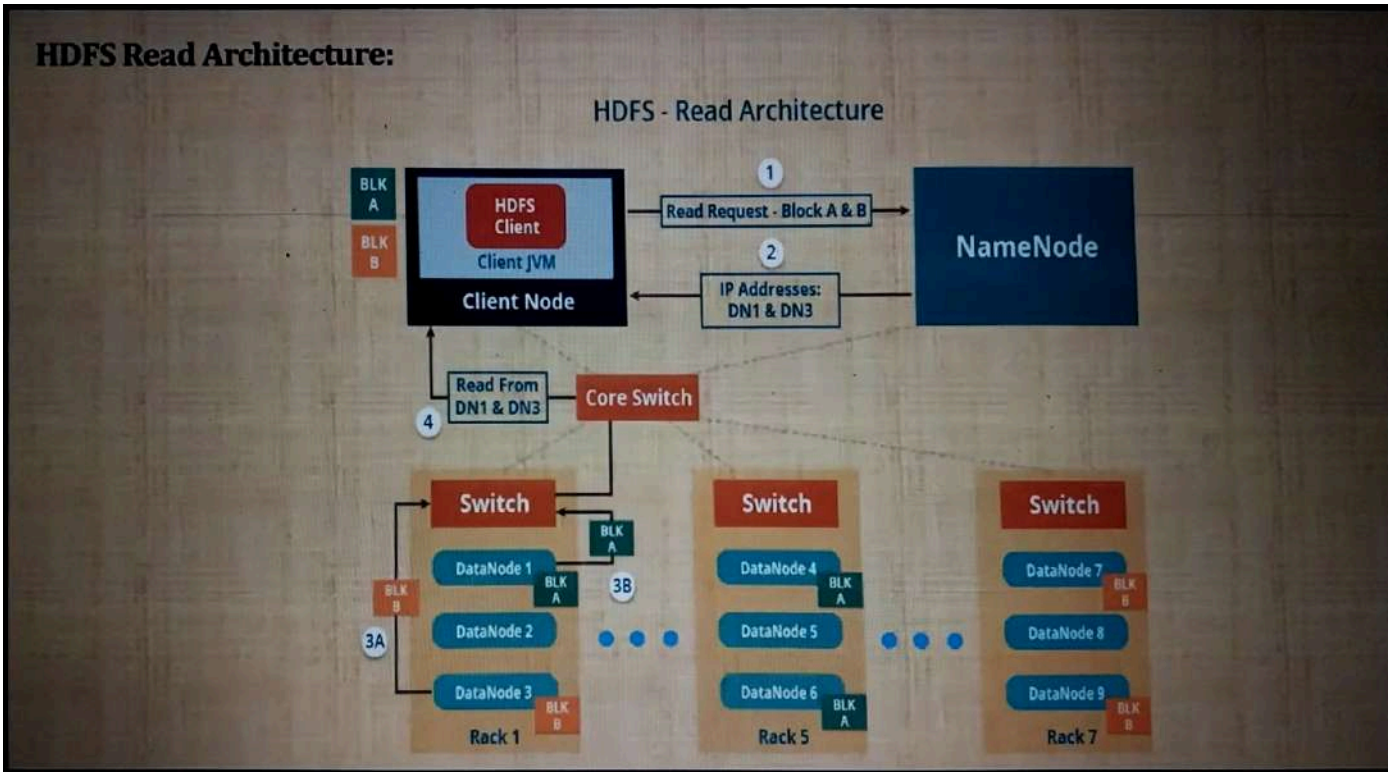

Advantages of MapReduce

1. Parallel Processing:

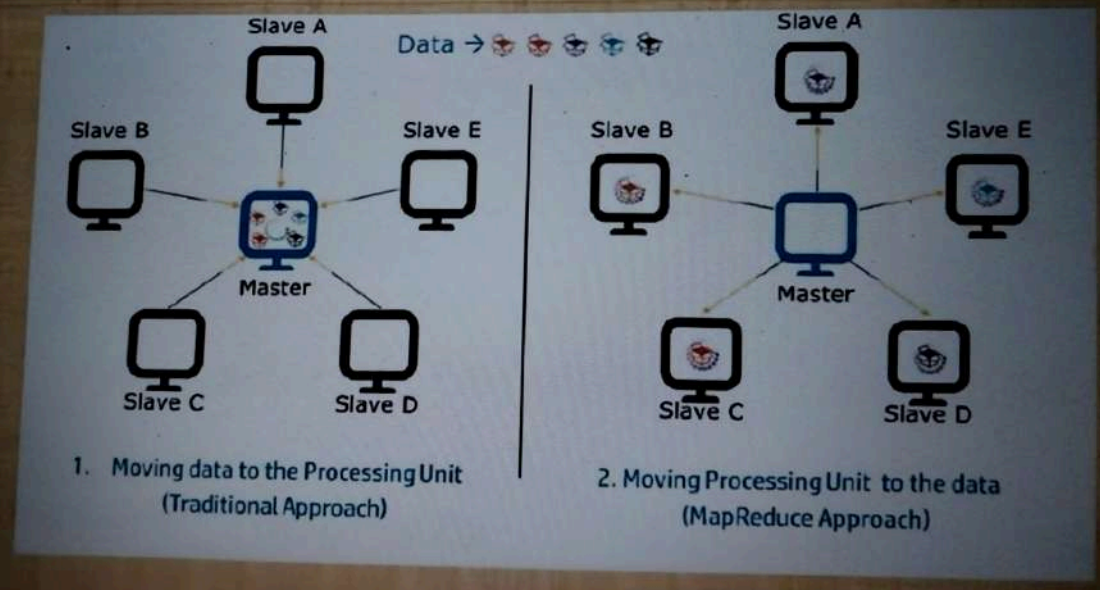

Map Reduce :-

T.A (Tradilathatal approch)

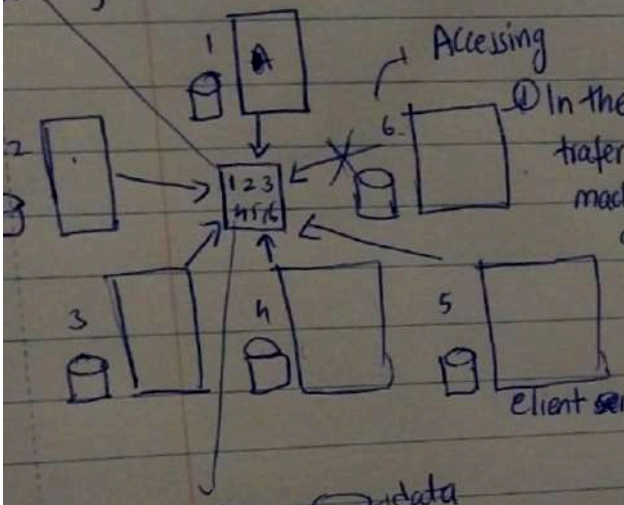

① In the processing of trafer the data one machine to another machine anthing can happenned.

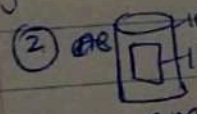

② one single server process the entire data.

B.D.A (Big Data approch)

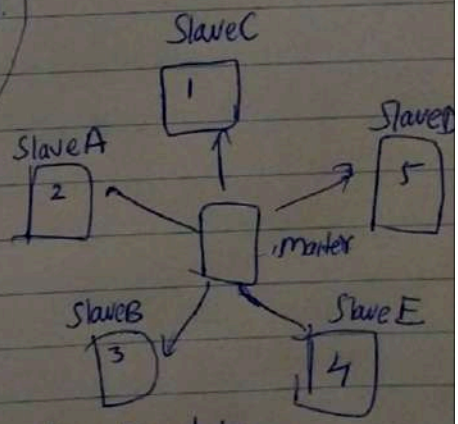

① here the data won't tranfer to one machine to another.

② The processing will take care of each slave.

③ It will only share the results of the slave's

0230214-020216_Recording_1440x900.mp4

Hadoop core components or daemons:-

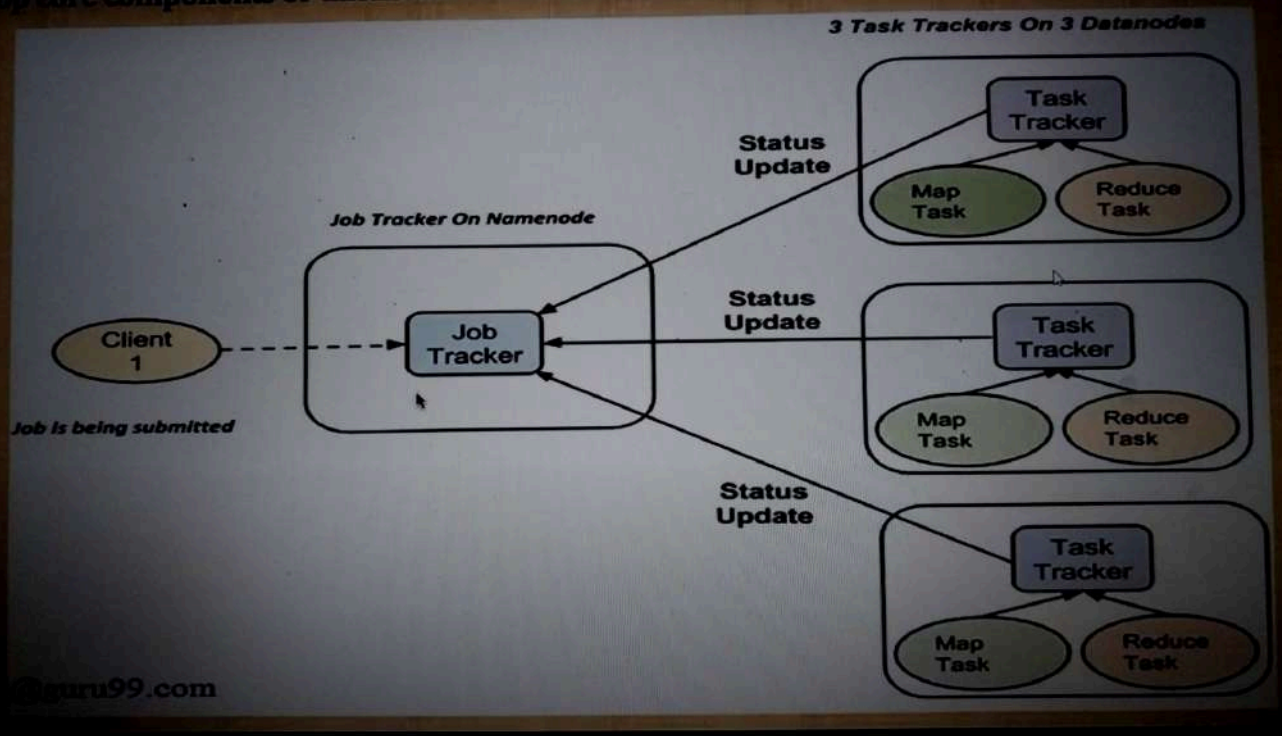

@guru99.com

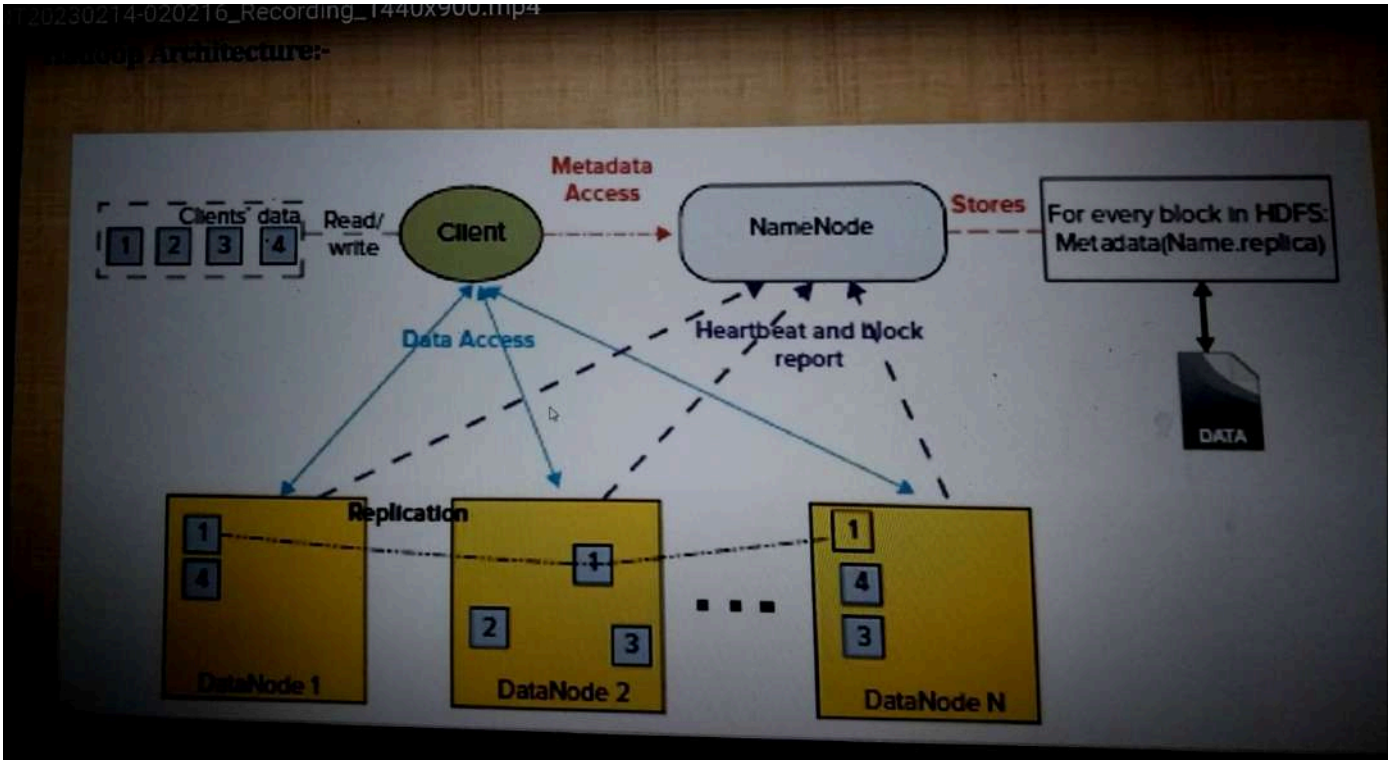
Metadata Access
Clients' data
1 2 3 4
Read/ write
Client
NameNode
Stores
For every block in HDFS: Metadata(Name.replica)
Data Access
Heartbeat and block report
DATA
Replication
1 4
1 2 3
1 4 3
DataNode 1
DataNode 2
DataNode N

Hadoop 1.X Limitation

Main drawback of Hadoop 1.X is that MapReduce Component in it's Architecture.

- → If the Name node server is down that total will down
- → It not suitable for Real-time Data Processing / Data Streaming
- → Job Tracker is the single point of failure
- → it runs only Map / Reduce jobs

Hadoop 2.X

YARN (Yet Another Resource Negotiator)

- Hadoop 2.X Allows to work in MR as well as other distributed Computing models like spark, Hama, Giraph, Message Passi
- 2.X Has better scalability, Scalable up to 10000 nodes per clu
- Multiple Namenode servers manage multiple namespace
- Can serve as a platform for a wide variety of data analytics - possible to run event processing, streming and real-time operations.

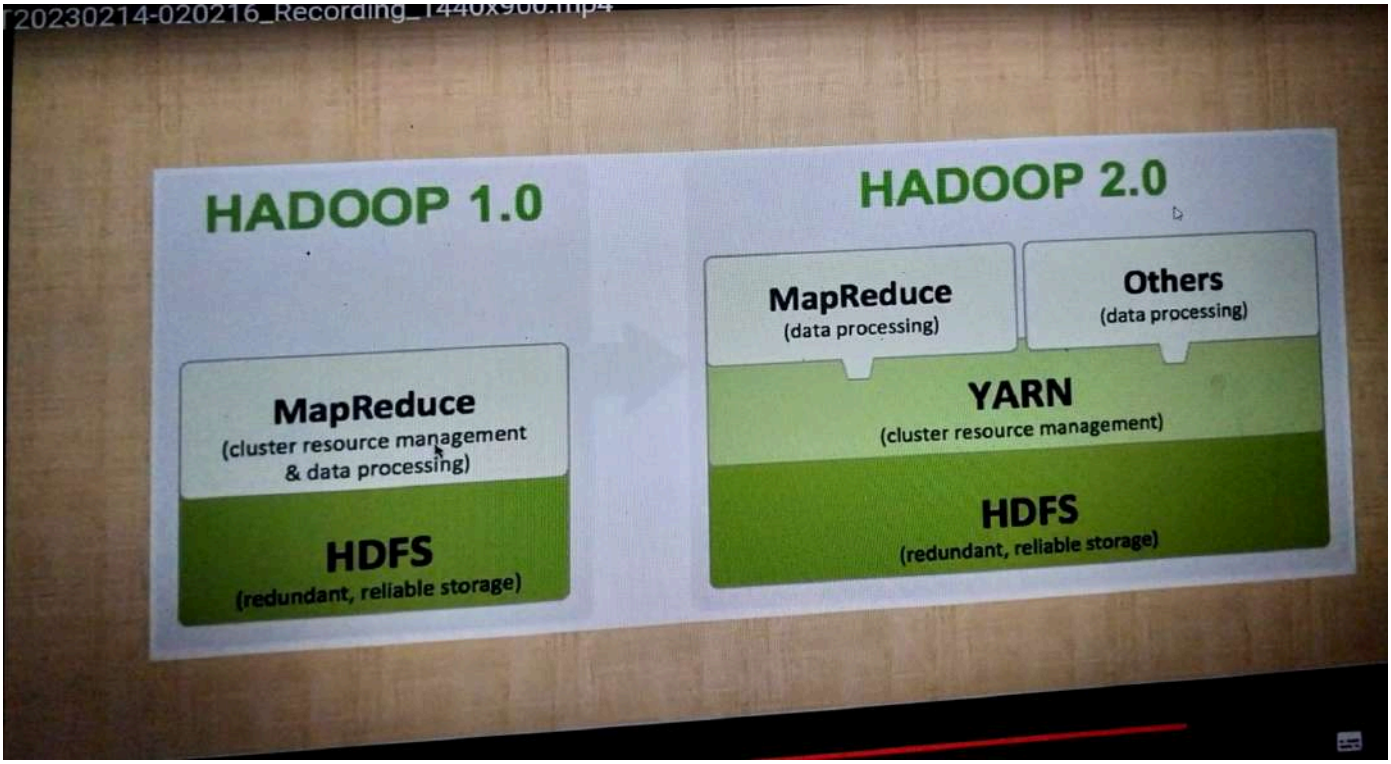
20230214-020216_Recording_1440x900.mp4
HADOOP 1.0
MapReduce
(cluster resource management & data processing)
HDFS
(redundant, reliable storage)
HADOOP 2.0
MapReduce
(data processing)
Others
(data processing)
YARN
(cluster resource management)
HDFS
(redundant, reliable storage)

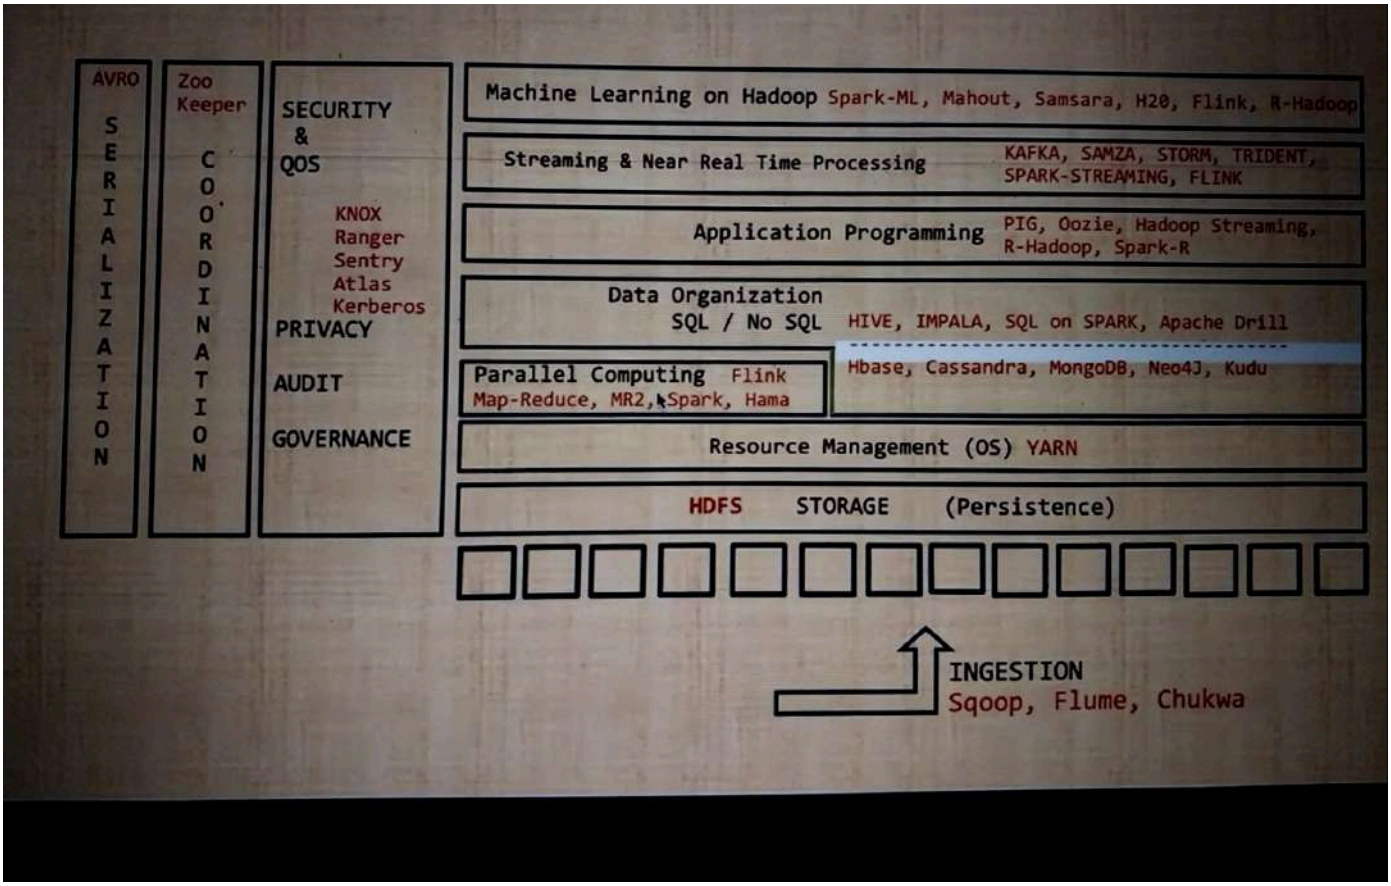
AVRO
SERIALIZATION
Zoo Keeper
COORDINATION
SECURITY & QOS
KNOX
Ranger
Sentry
Atlas
Kerberos
PRIVACY
AUDIT
GOVERNANCE
Machine Learning on Hadoop Spark-ML, Mahout, Samsara, H20, Flink, R-Hadoop
Streaming & Near Real Time Processing
KAFKA, SAMZA, STORM, TRIDENT, SPARK-STREAMING, FLINK
Application Programming
PIG, Oozie, Hadoop Streaming, R-Hadoop, Spark-R
Data Organization
SQL / No SQL
HIVE, IMPALA, SQL on SPARK, Apache Drill
Parallel Computing Flink
Map-Reduce, MR2, Spark, Hama
Hbase, Cassandra, MongoDB, Neo4J, Kudu
Resource Management (OS) YARN
HDFS STORAGE (Persistence)
INGESTION
Sqoop, Flume, Chukwa

-1.X — Job tracker

↳ Resource manangment & processing

2.X.V YARN SPARK

Hadoop 2.X Version

↳ YARN → Resource management

↳ SPARK → Processing

YARN → Yet Another Resource Negotion

→ Container Concepts

↳ memory / duck / cpu cycles

Job: app master

1:1 sure.

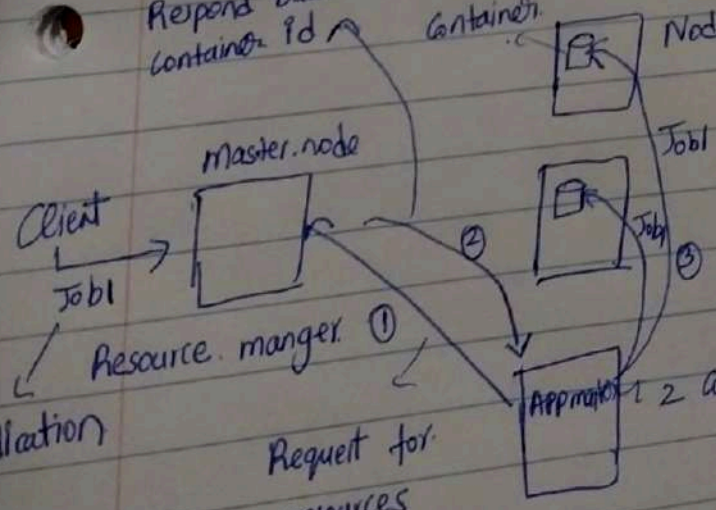

Resoure manager.

① It create a container

② It create app master

The node manager will give heart beat to the Resource manger. health of each Node is updated to node manger to. resource ma

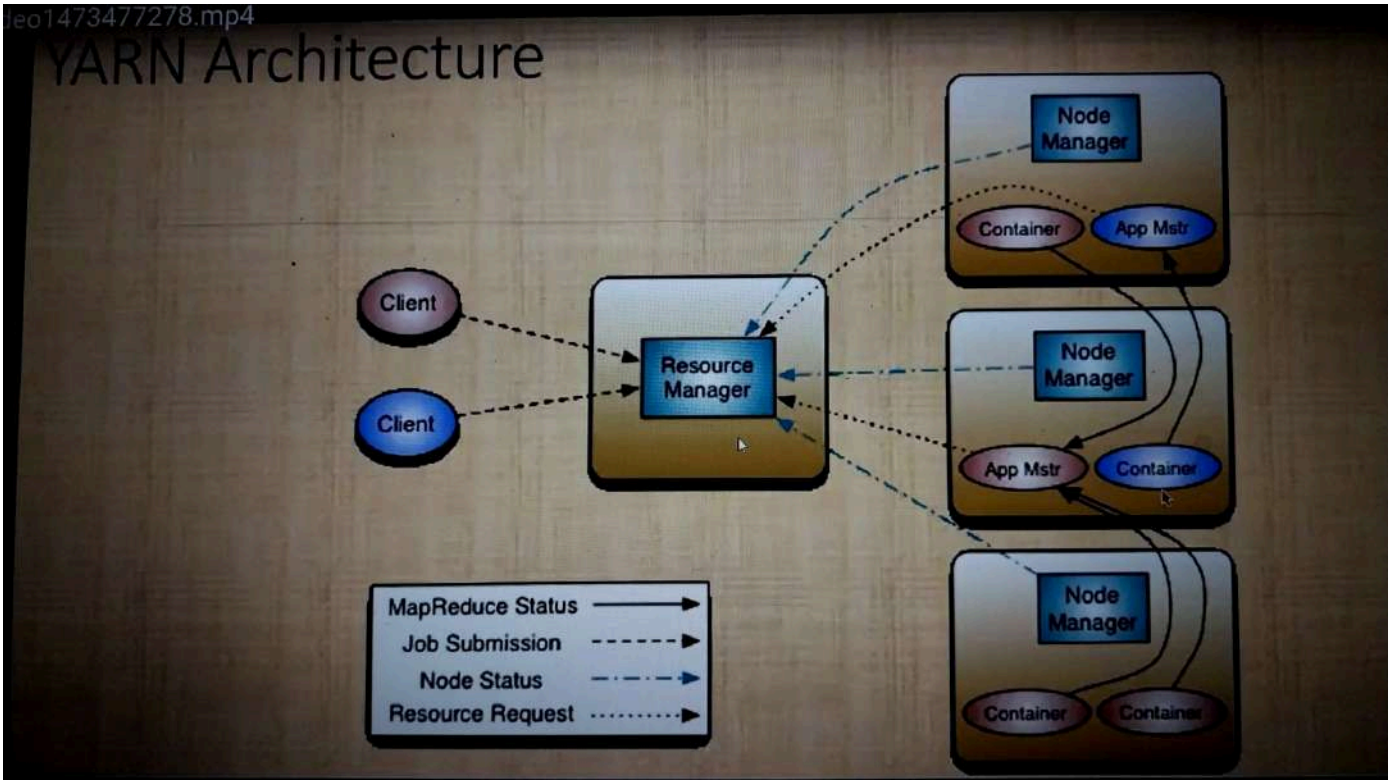
eo1473477278.mp4
YARN Architecture
Node Manager
Container
App Mstr
Client
Client
Resource Manager
Node Manager
App Mstr
Container
Node Manager
Container
Container
MapReduce Status
Job Submission
Node Status
Resource Request

Running a YARN application

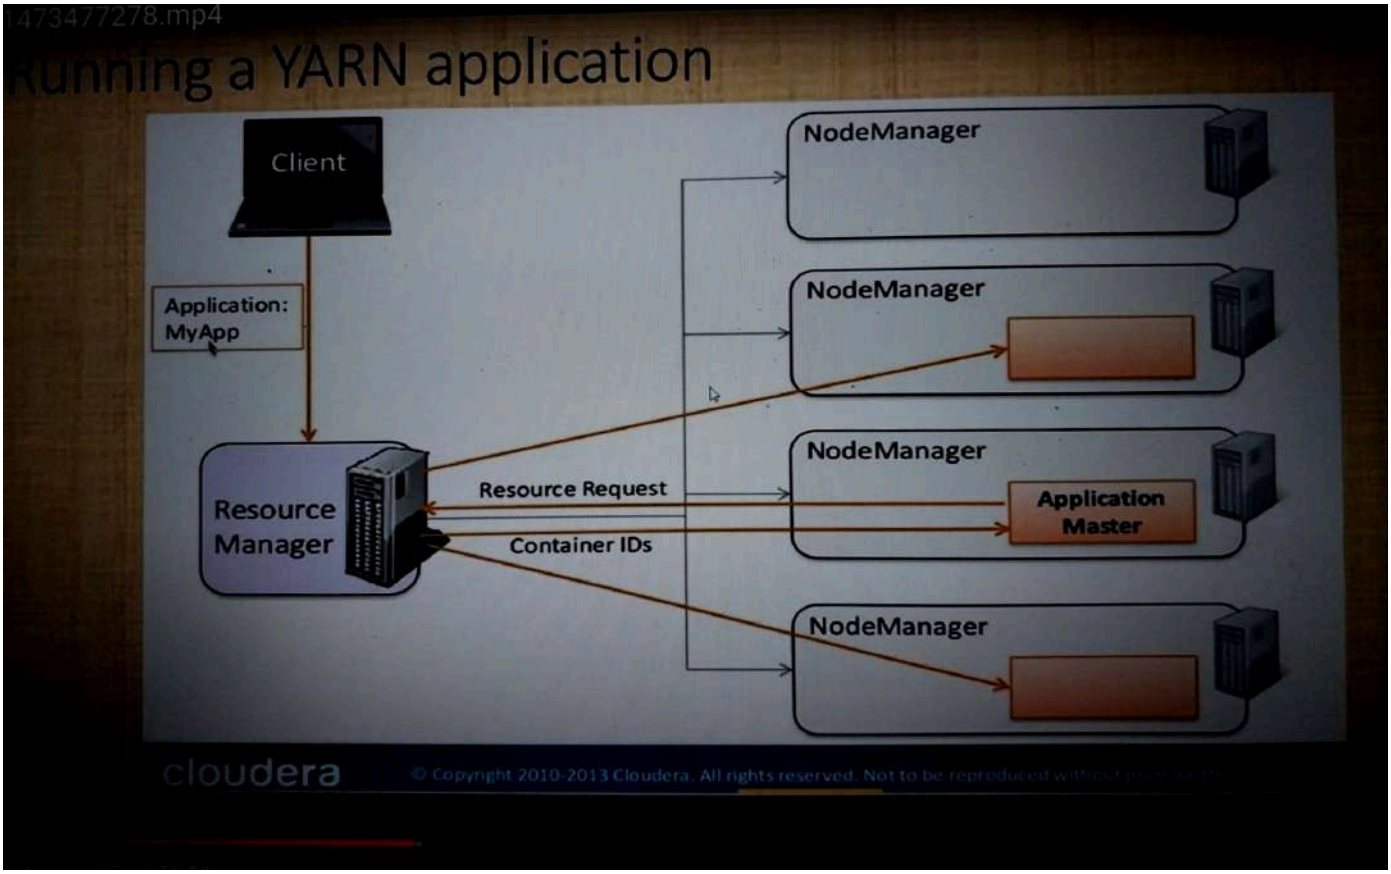

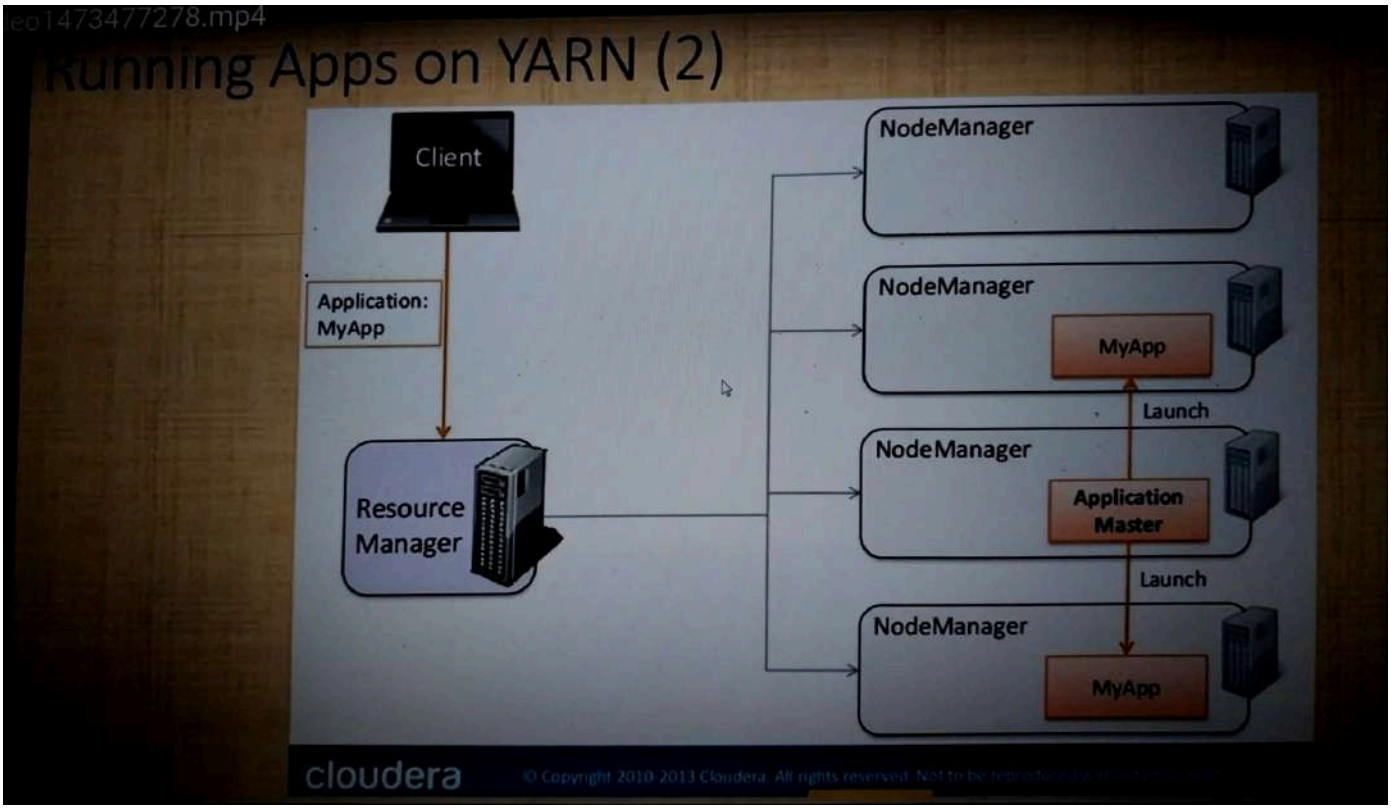
eo1473477278.mp4
Running Apps on YARN (2)
Client
Application: MyApp
Resource Manager
NodeManager
NodeManager
MyApp
Launch
NodeManager
Application Master
Launch
NodeManager
MyApp
cloudera
© Copyright 2010-2013 Cloudera. All rights reserved. Not to be

CapEx vs OpEx

Capital Expenditure	Operational Expenditure
Spending on infrastructure is completed upfront	No upfront cost
Cost written off over a period of time	Pay for service as you consume it
	Deduct from tax bill in same year as expense occurs

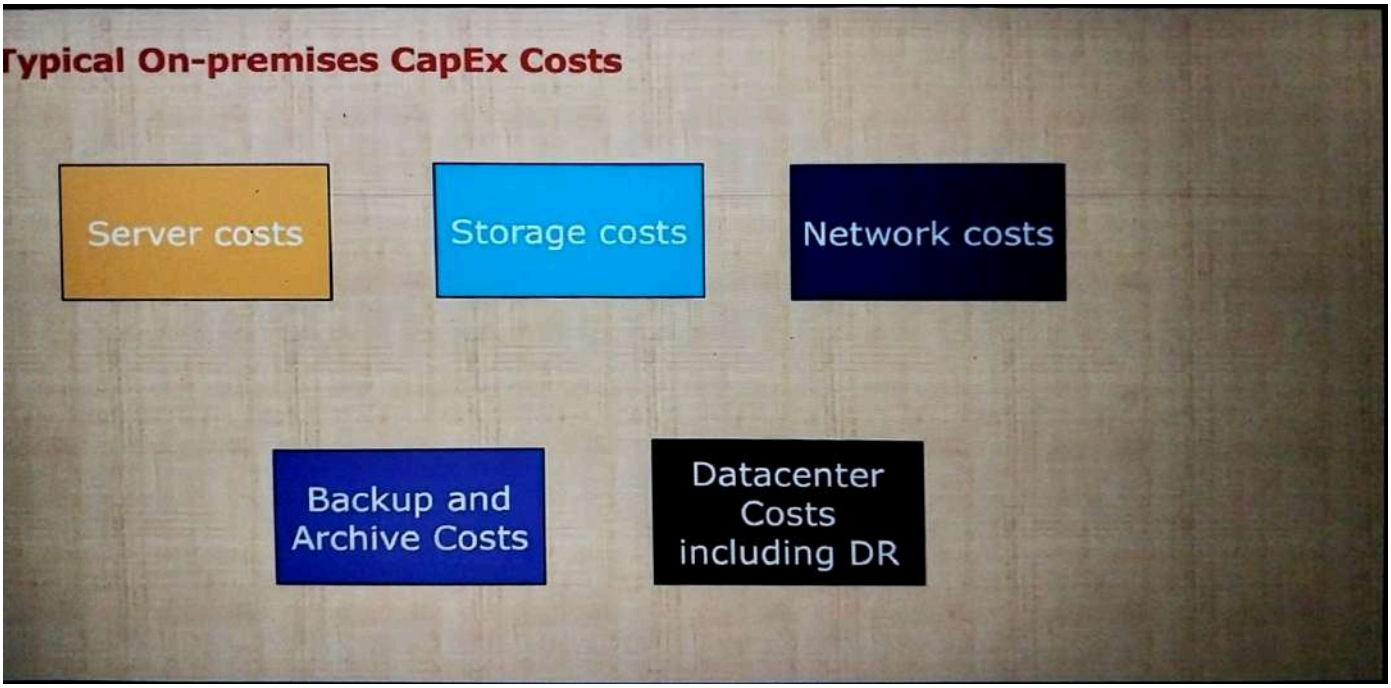
Typical On-premises CapEx Costs
Server costs
Storage costs
Network costs
Backup and Archive Costs
Datacenter Costs including DR

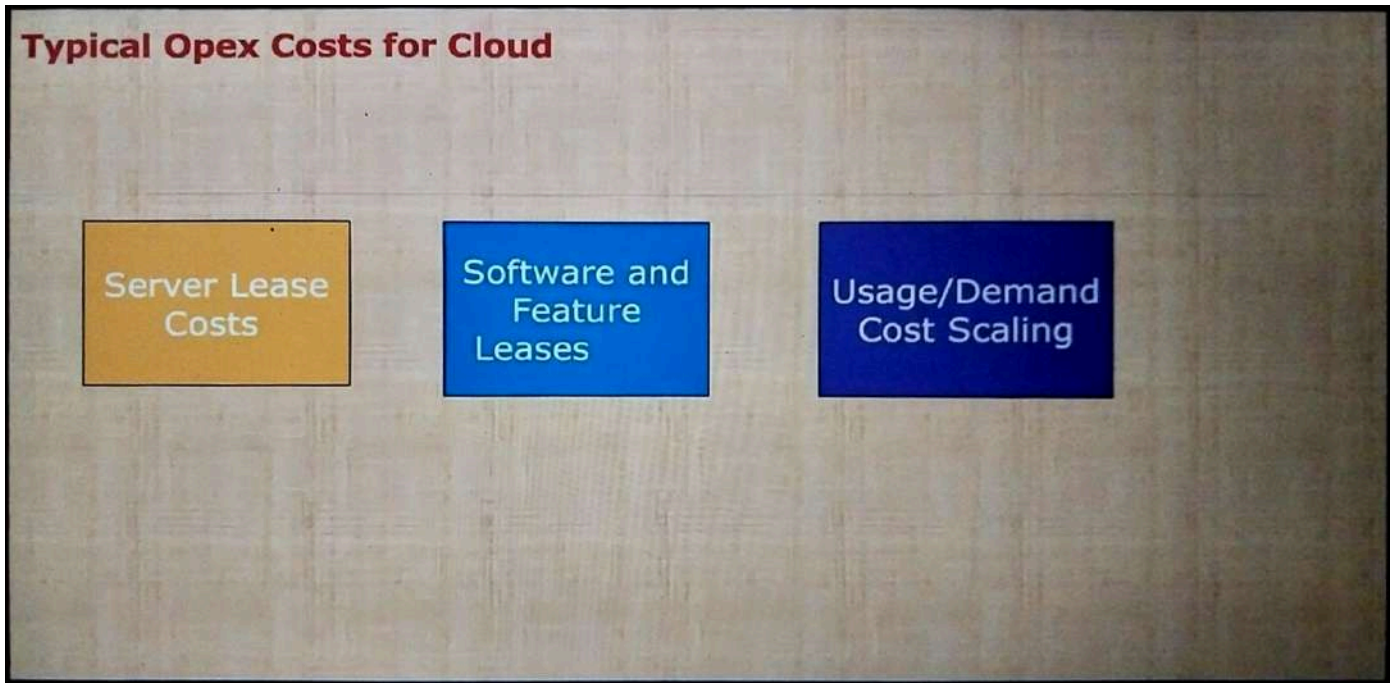
Typical Opex Costs for Cloud
Server Lease Costs
Software and Feature Leases
Usage/Demand Cost Scaling

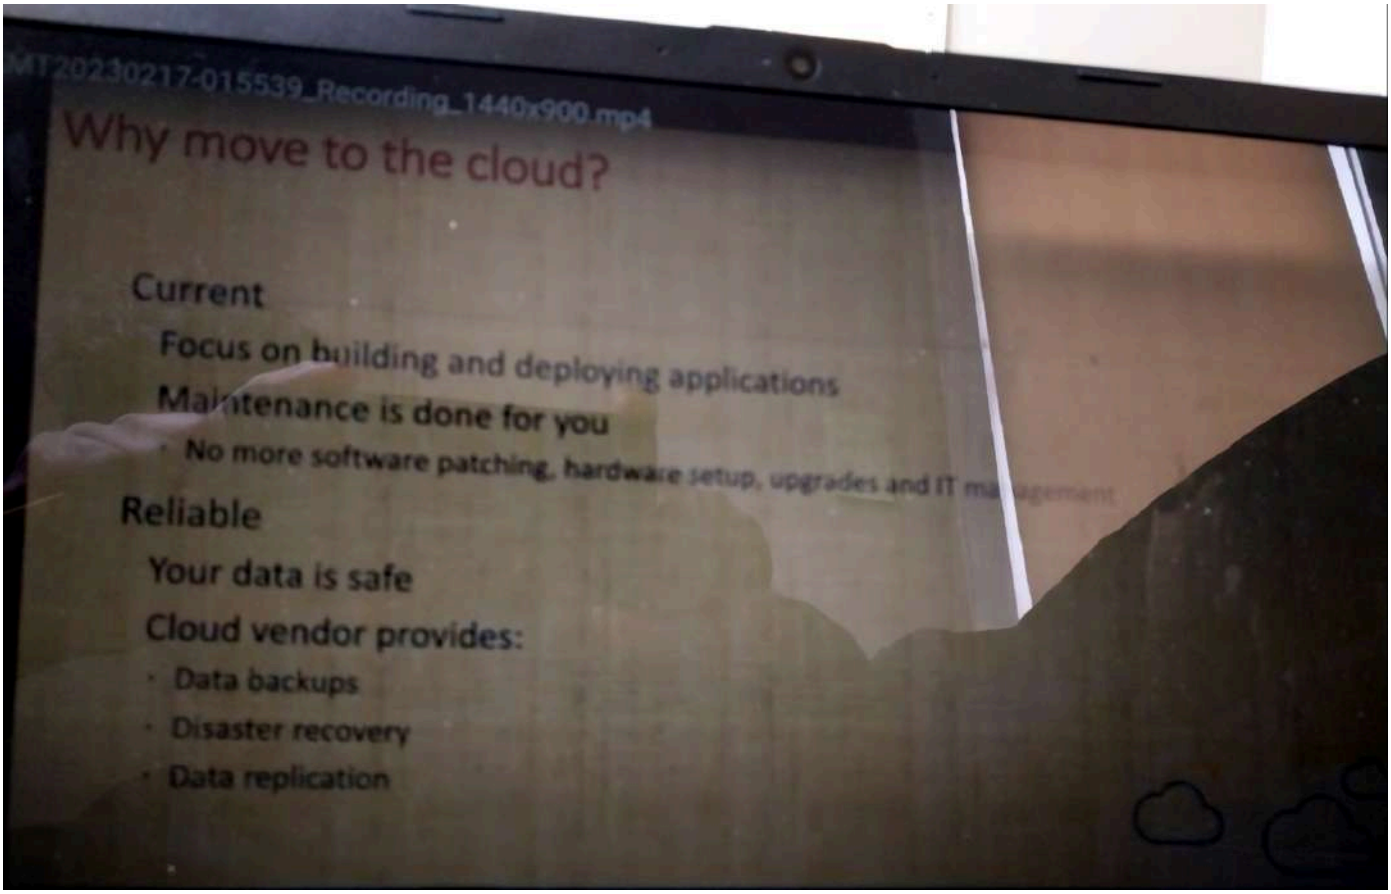
MT20230217-015539_Recording_1440x900.mp4
Why move to the cloud?
Current
Focus on building and deploying applications
Maintenance is done for you
No more software patching, hardware setup, upgrades and IT ma
Reliable
Your data is safe
Cloud vendor provides:
Data backups
Disaster recovery
Data replication

Spark → Parallel processing platforms

(In-memory computation)

M.R (Map Reduce)

→ 2 Functions map & Reduce
→ batch processing
→ Disk to process the data

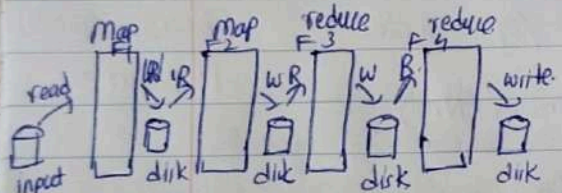

→ Multiple Read/write operation
→ Time - processing is more/high

↳ on Disk takes more time

RAM → Memory → we can reduce time

ROM → DISK → Time is huge

Spark

→ multiple Fuctions
→ both batch & real-time data
→ Memory - process the data

Spark (for processing)

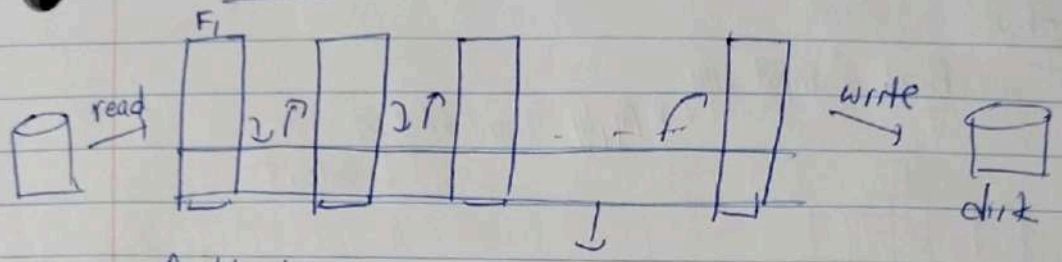

Read/write happens in the RAM memory

In memory / distributed memory

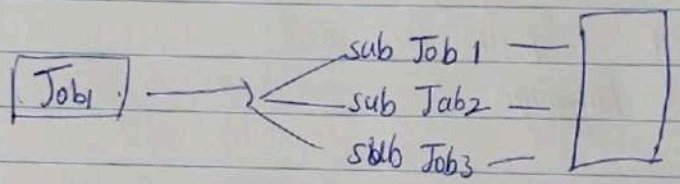

Spark

2 operation

① . Trasformation — Multiple Fuction. (Converting one form of data into another form)

② . Action → Multiple. Fuction
↳ show the result

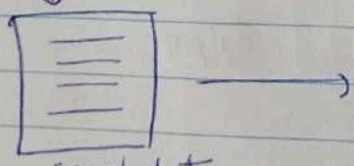

sample.txt.

This is a BigData session.

→ ["This", 'is', 'a', 'BigData', 'session']

Python

Read a input file
↳ Apply Transformation
↳ output.

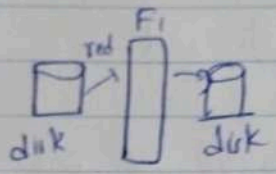

Spark + Python — written Scala Language
↳ pyspark ⇒ python language
↳ language

Pyspark

lazy Evaluation
↳ Transformation - unique feature
! will not consume any cpu cycle / ~~execute~~ not execute until we call action

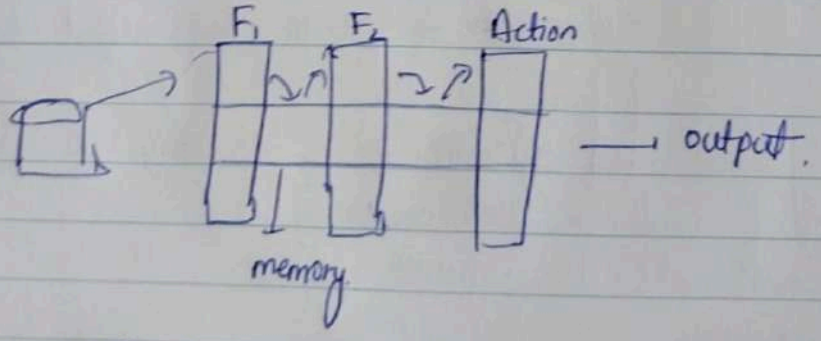

Transf
- map()
- filter()
- ReduceByKey()

Action
- collect()
- show()
- take()

} shows output

spark → API's

spark → 3X.V

1.4 ≤ 1.4 >

↓ ↓

RDD API DF ← RDD

&

SQL

RDD → unstructured Data — output [.....]

↓

DF → structured way of processing Data —

& SQL

└ Hand-on - cloud

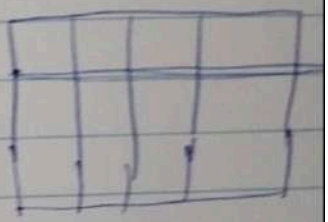

	Master	Slave
HDFS -	Name Node	Data Node
MR -	Job Traker	Task Traker
YARN -	Resource manager	Node manger
SPARK -	Drivers	Executors

Use case - cloud ? why I should move to cloud

start up - in developing a mobile App - Food ordering App

End - End product Development

client → Development → Deliver (customer)

launch it

Infrastructure : → server
(server, computers) OS
storage
software stack ← tool functality

↓

Avg Computing power

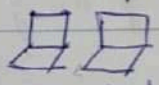

2 Laptops

10-developer's

Bengluru - 10,000 customer

→ Infrastructure for production (launch it) is different compare to development.

Day1 Target customer 10,000

Intra required

1,000 5,000

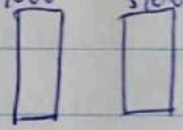

server1 server2

IT Team 2 members

No. of customer - 2k - 5k

Day5 - 9K customer → works fine

Day6 - 15K customer → Accell the App

? - server down

- customer - bole & bad Feedback

Add New servers → 2 server - Budget
- configure

Day 10 — 20,000 — 4 server's — 3 IT Team.
No issues ✓

Day 20 → 30,000 — 6 server - 4 IT Team
↓ today - No. issues ✓

Day 30.

Day 32 - No. of user's - 5K
33 " - 6K
" " - 10K
? Cost is fixed ?

observation

Focus ? Developing New product with less budget.

- → maintance issues - costly, Time
- → server cost
- → IT Team.
- → cost is fixed

3rd party — comes to you
(They take of all Infastructure (serve's) Accepted

after — cloud.
pay for what you use.
Day - 4 hrs - pay for 4 hrs
1 hr → " "

Where can infrastructure be hosted?
- For an enteprise, it is on a Data Centers

What is cloud Computing
- Delivery of computing servicei over internet

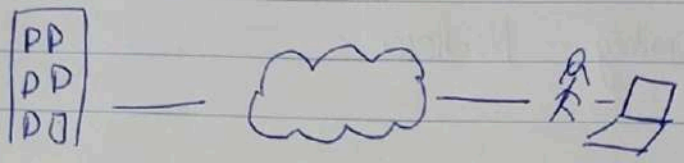

-) accessing someone elie ~~over~~ computer over the internet for my computer and pay only for cloud servicei you uie.

What is Cloud Computing?
* Renting resourcei vs Purchasing the hardware
* Pay for what you uie
* Run your applications in someone else's datacenter
* Cloud provider is reiponsible for the physical hardware and facilitiei necessary to execute your work
* Cloud provider is reiponsible for keeping the seriei to uptod

Elaiticity)

→ As your workload changes, resources can be changed to compensate (up or down)

example : Seasonal Demand for retail website Black Friday

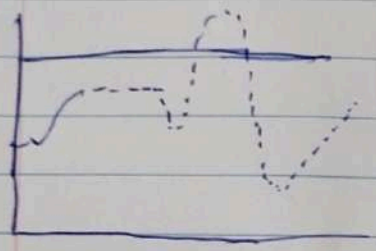

static

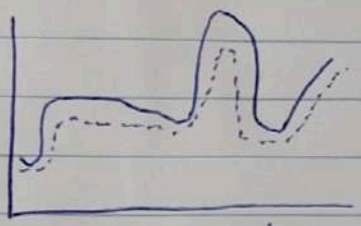

Elastic Scaling

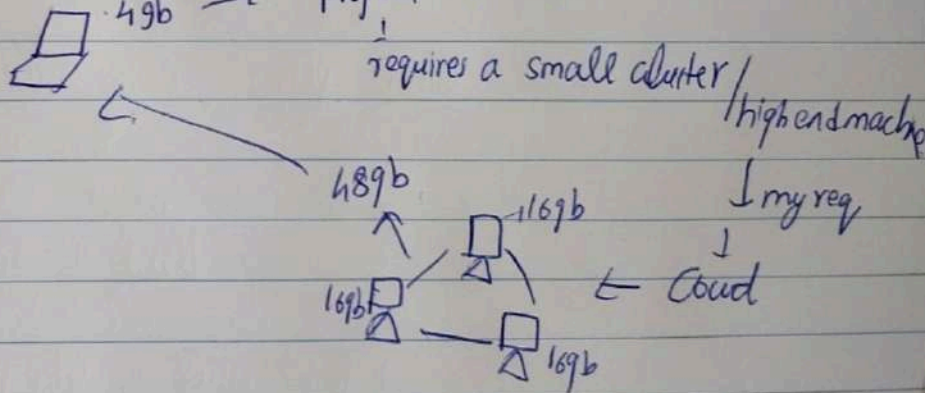

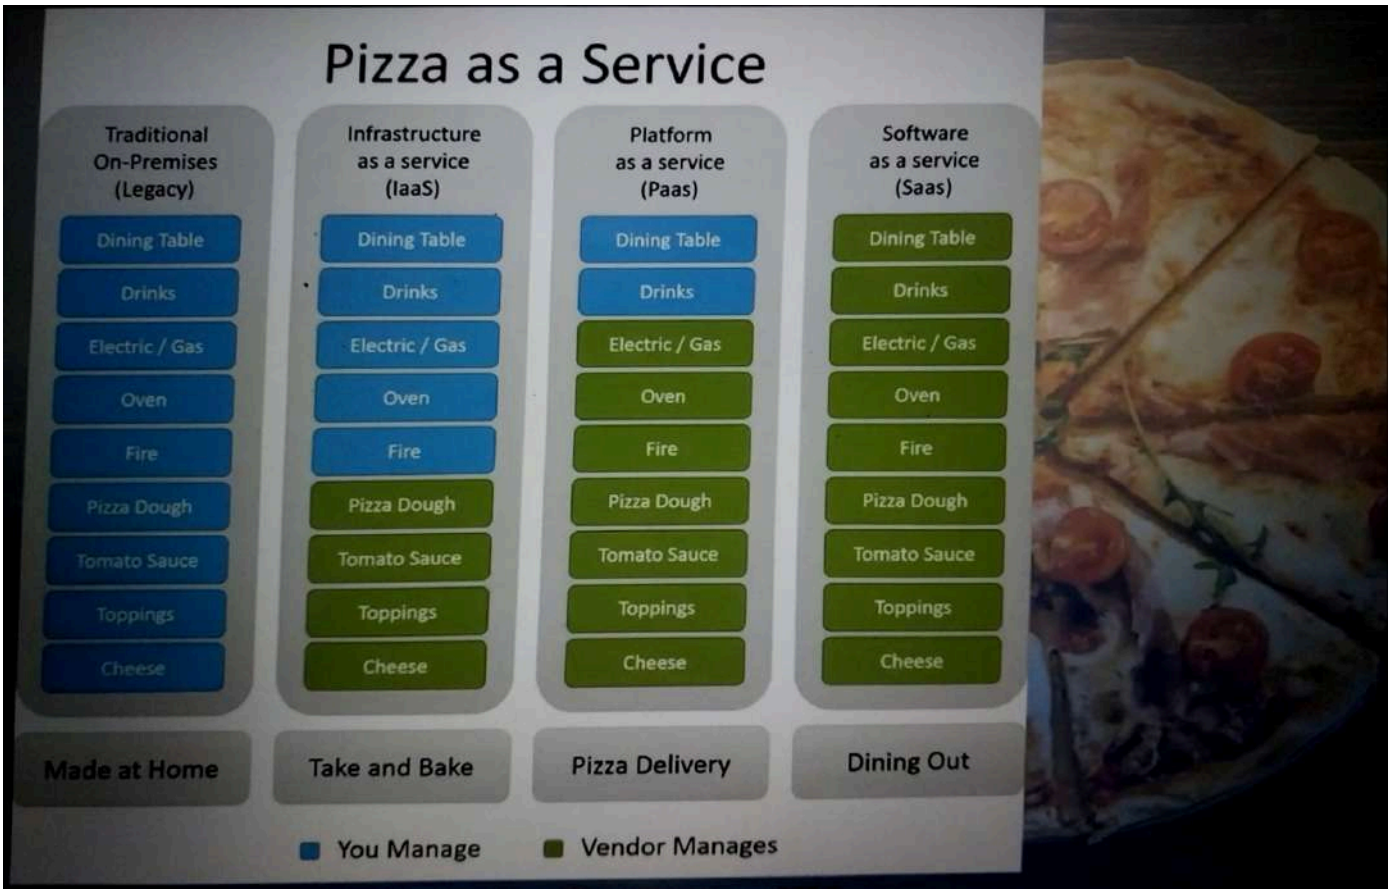
Pizza as a Service
Traditional On-Premises (Legacy)
Infrastructure as a service (IaaS)
Platform as a service (Paas)
Software as a service (Saas)
Dining Table
Drinks
Electric / Gas
Oven
Fire
Pizza Dough
Tomato Sauce
Toppings
Cheese
Made at Home
Take and Bake
Pizza Delivery
Dining Out
You Manage
Vendor Manages

XaaS (What can be rented?)

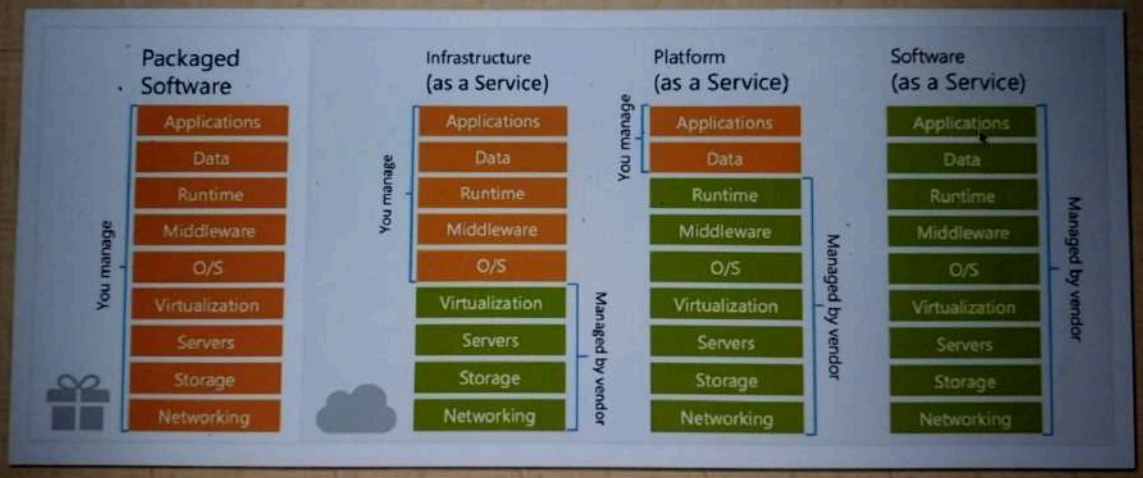

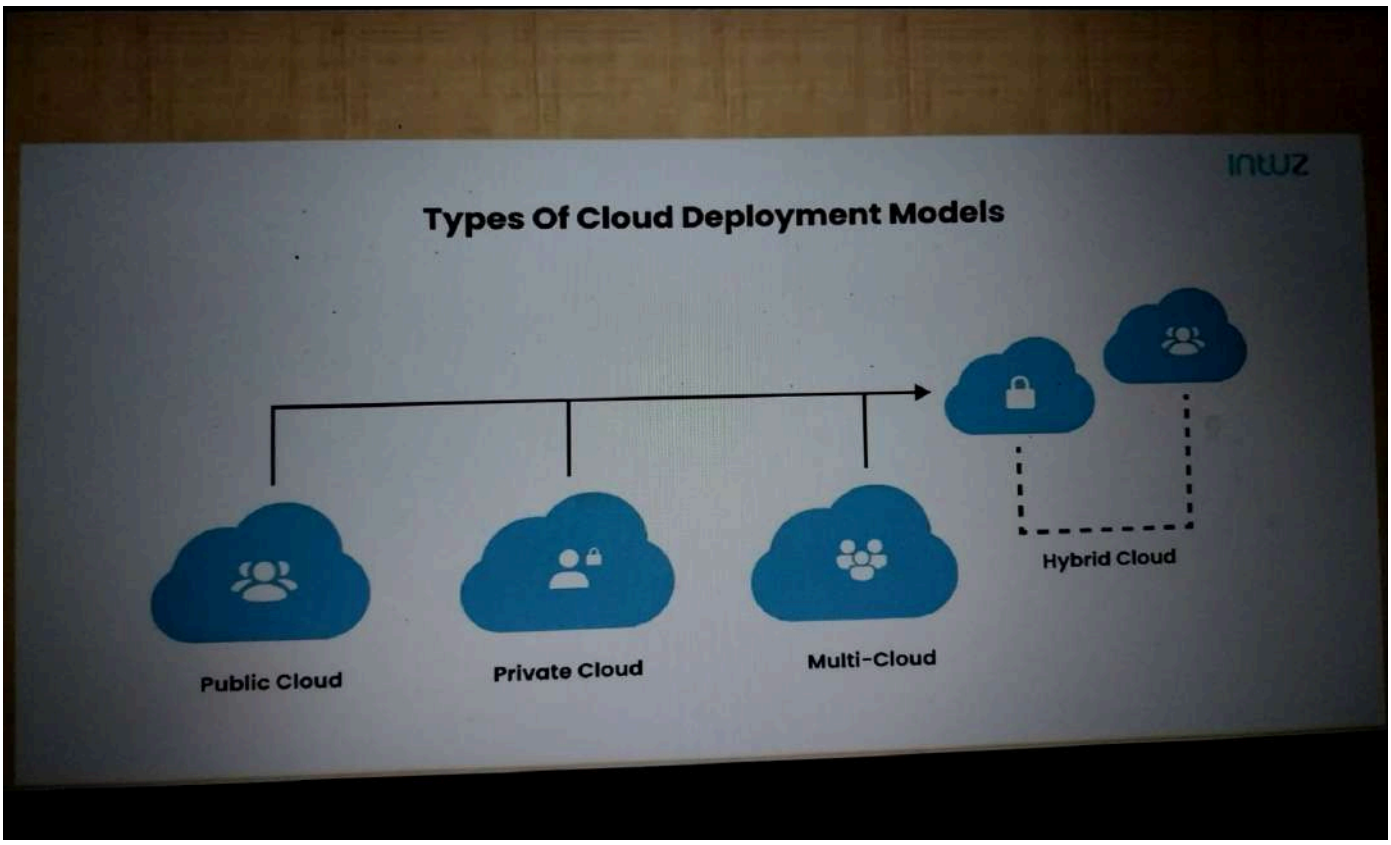
Types Of Cloud Deployment Models
Public Cloud
Private Cloud
Multi-Cloud
Hybrid Cloud

What is Azure?

- A Cloud Computing Platform from Microsoft
- Released as **Windows Azure** in February 2010
  Renamed to **Microsoft Azure** on March 25, 2014
- Provides a web portal to access and manage cloud services & resources.
- Free to start, pay-per-use

Data Centers

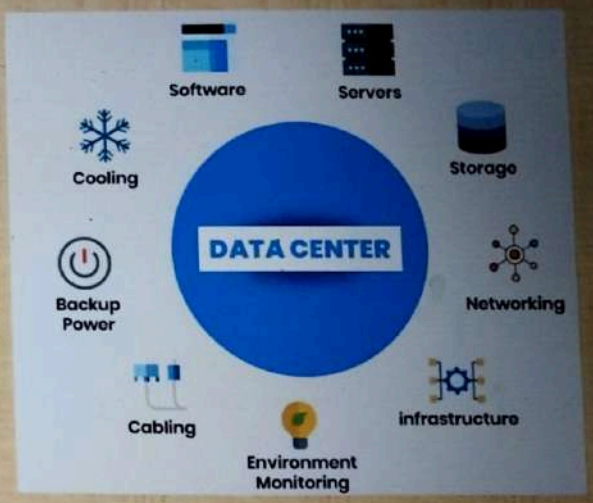

https://news.microsoft.com/innovation-stories/project-natick-underwater-datacenter/

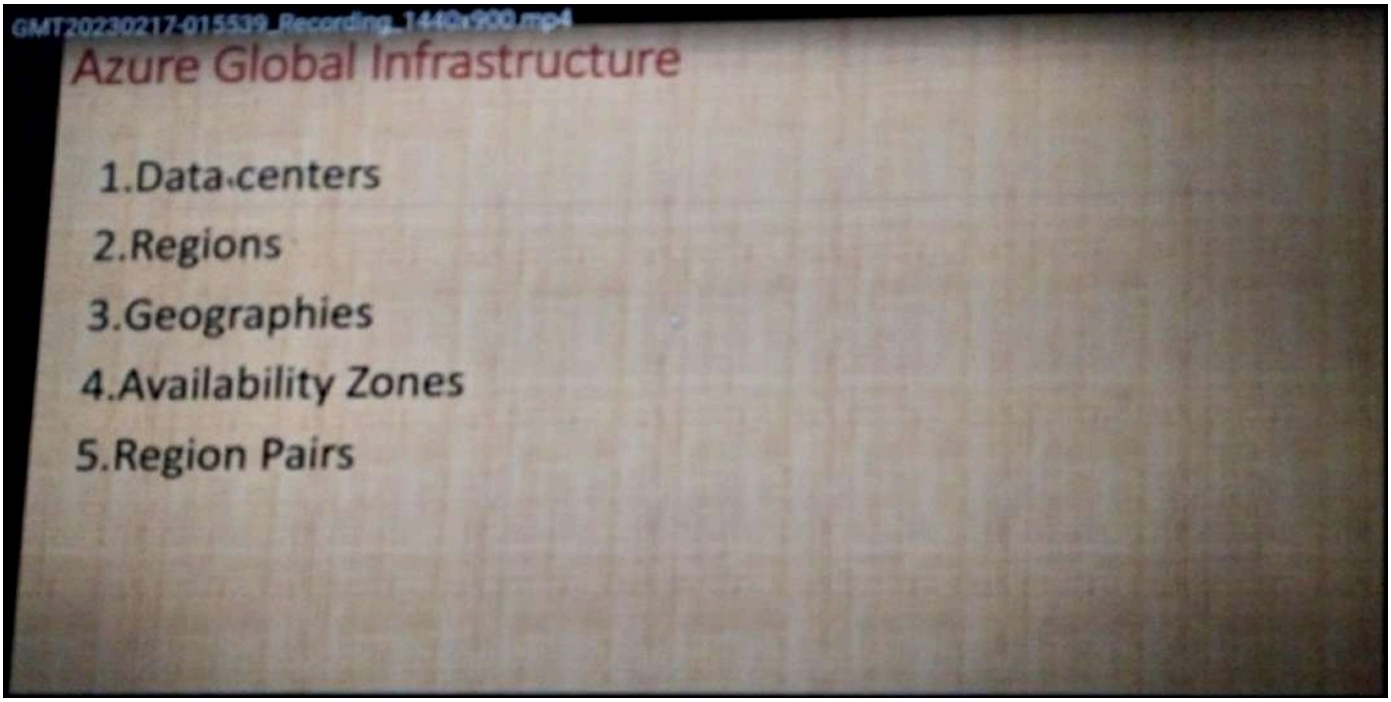
GMT20230217-015539_Recording_1440x900.mp4
Azure Global Infrastructure
1.Data centers
2.Regions
3.Geographies
4.Availability Zones
5.Region Pairs

Azure Regions

Location for your resources

Area containing at least one datacenter

Select a region when deploying a resource

Azure Geographies

An Azure geography is an area of the world that contains at least one Azure region.

Ex: United States, United Kingdom, India, Asia Pacific etc

Azure Government

- This geography is only available to the United States federal, state, local, and tribal governments and their partners.

Azure availability zones

Availability zones are physically separate datacenters within an Azure region. Each availability zone is made up of one or more datacenters equipped with independent power, cooling, and networking. An availability zone is set up to be an *isolation boundary*. If one zone goes down, the other continues working. Availability zones are connected through high-speed, private fiber-optic networks.

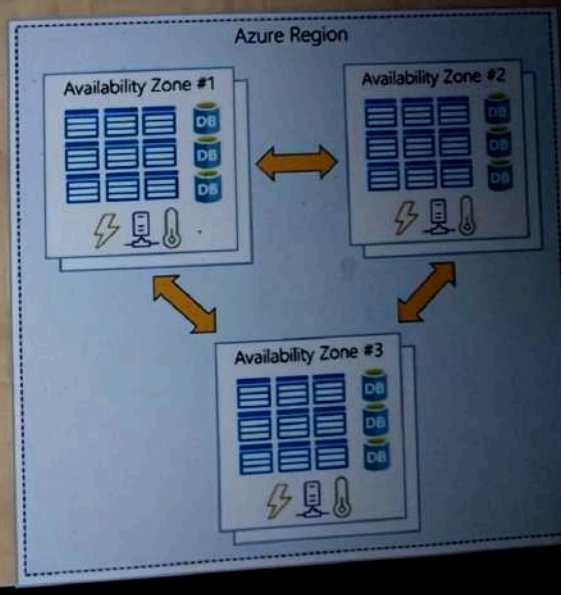

Each Azure region is always paired with another region within the same geography (such as US, Europe, or Asia) at least 300 miles away. This approach allows for the replication of resources (such as VM storage) across a geography that helps reduce the likelihood of interruptions because of events such as natural disasters, civil unrest, power outages, or physical network outages that affect both regions at once. If a region in a pair was affected by a natural disaster, for instance, services would automatically failover to the other region in its region pair.

Examples of region pairs in Azure are West US paired with East US and SouthEast Asia paired with East Asia.

Azure region pairs

Geography	Regional Pair A	Regional Pair B
Canada	Canada Central	Canada East
China	China North	China East
India	Central India	South India
Japan	Japan East	Japan West
North America	East US	West US

Core cloud services

Compute Storage Networking
App Services Analytics

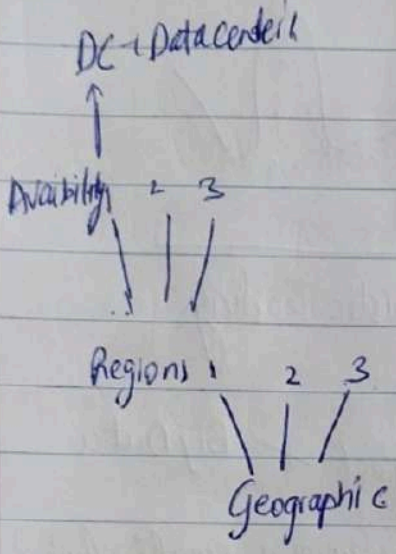

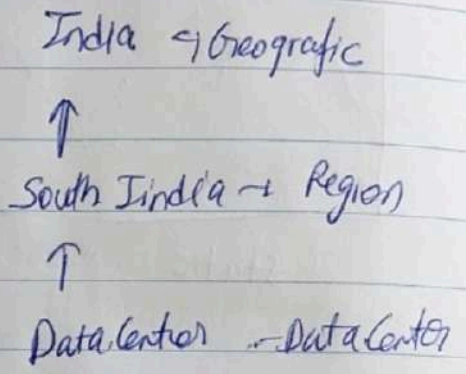

IAAS
- → Compute Engine
- → Storage services
- → Network services

Compute engin
- → VM
- → Containers
- → Kubernetes

Storage servic
- Structure data → Azure SQL
- Semi-Structure data → Cosmos DB
- un-structured data → ADL & Blob

network servi
- → NIC
- → IP address
- → subnets

Resource Group

R.G .....
- VM
- storage

Service — Azure

* Storage services

store any data

- Structured Data → Azure
- Unstructured Data → Blob - ADLS
- Semi-structured Data → or COSMOS DB

col / attributes

rows / tuples

schema.

- No file system
- Flat file

Relational Data

Schema on write → RDBMS

1) Create the schema before / load / store data

Schema on read — non-relational data.

→ NoSQL — Not only SQL

can store SQL data and also

→ flexibility in schema

Structured data ⟶ semi-structed data.

①

create table student1 (Sid number, name varchar(20) ....);

Mango.DB → NoSQL → (JSON Java Script object Notation)

id	Name	H1	H2	H3	....	H10
1	Anu	-	-	2	?	
2	-	-	-	?	?	-
3	-	-	-		?	?

→ only have H3 all are empty so waste of storage

Limitation are in the back

1) Every row/record in a table have same no. of attribute

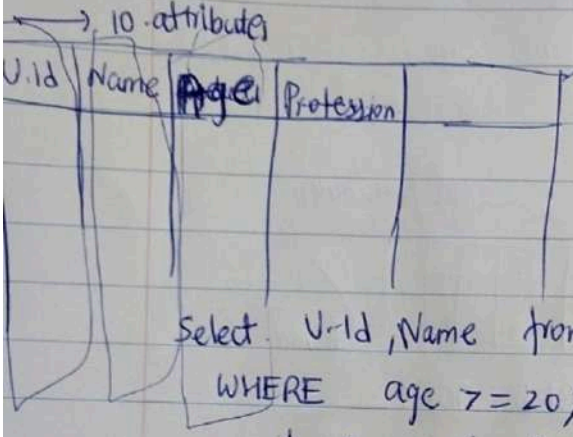

Select U-Id, Name from Userdetails
WHERE age >= 20;

it process in the Traditinal way (because it check all attributes it was not necessary).

② SQL is not an efficient for reading

NOSQL
Mango-DB
→ JSON
Java script object Notation

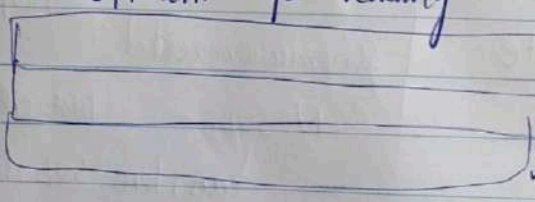

↓ it change column to rows
Columnar Data store

```
{
 U-id:
 Name:
 hobby: [" ", " ", " "]
}
```

```
{
 U-id:
 Nam:
 hobby: [ ..... ];
}
```

↳ we can take how many you want

Bigdata
↳ processing, ADLS → Big Data processing required to have file system

Flat Files
=
Blob

Storage Account → ☒ to enable
* we need not to enable
to hir namespace

↑
container - space
↑
store our data

BigData processing
=
ADLS

→ Storage
☑ to enable the hir.
namespace
↑
container
↑
Can create a diretory fil
↑
store our data

Access
Hot - Use more
cool - Not use more
Archive -

Storage services

Structured — Azure SQL → DB → SQL server
Semi structured — Cosmos DB → Mango DB → JSON Node
un-structured → blob & ADLS

Azure Databricks

Big data → Hands on

Spark → run all spark job — Introduce the service Name Azure Databricks
→ programing Language — Scala — pyspark
→ platform - spark - 3. X. version
→ service - data. bricks

Azure Databricks

1) Create a Databricks workspace
2) Create a cluster — single Node Cluster
3) Create a Notebook — write spyspark code
4) DBFS → Data bricks file system
   L create DBFS & storege datasets

Spark

Version ≤ 1.4 → Spark RDD

Version > 1.4 → 3.X

Spark DF & SQL — object spark session → spark

Spark RDD — spark context → SC

Processing the data in the structured & semi structure

unstructured list of list [[ ]]

Pyspark → DF (col) / SQL

pyspark - programing + SQL - query language

rows

Spark

2 option

1) Transformation → it will not execute → a, b, c → Fuction → ['a', 'b', 'c']

2) Action

not easy memory

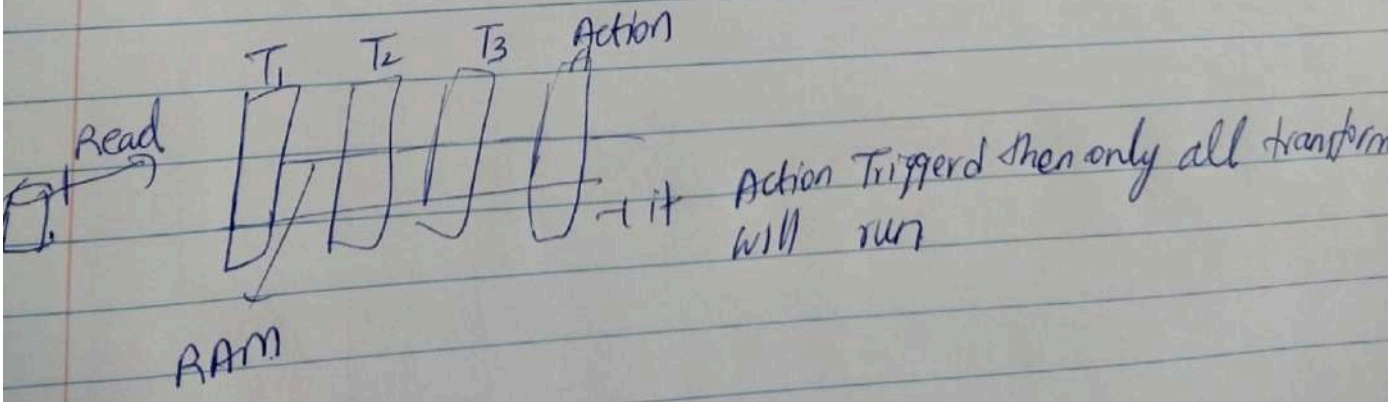

→ it Action Triggerd then only all transform will run